

HAR_PREDICTION_MODELS

March 16, 2019

```
In [1]: import numpy as np
import pandas as pd
```

0.1 Obtain the train and test data

```
In [2]: train = pd.read_csv('UCI_HAR_dataset/csv_files/train.csv')
test = pd.read_csv('UCI_HAR_dataset/csv_files/test.csv')
print(train.shape, test.shape)
```

(7352, 564) (2947, 564)

```
In [3]: train.head(3)
```

```
Out [3]:
```

	tBodyAccmeanX	tBodyAccmeanY	tBodyAccmeanZ	tBodyAccstdX	tBodyAccstdY	\
0	0.288585	-0.020294	-0.132905	-0.995279	-0.983111	
1	0.278419	-0.016411	-0.123520	-0.998245	-0.975300	
2	0.279653	-0.019467	-0.113462	-0.995380	-0.967187	

	tBodyAccstdZ	tBodyAccmadX	tBodyAccmadY	tBodyAccmadZ	tBodyAccmaxX	...	\
0	-0.913526	-0.995112	-0.983185	-0.923527	-0.934724	...	
1	-0.960322	-0.998807	-0.974914	-0.957686	-0.943068	...	
2	-0.978944	-0.996520	-0.963668	-0.977469	-0.938692	...	

	angletBodyAccMeangravity	angletBodyAccJerkMeangravityMean	\
0	-0.112754	0.030400	
1	0.053477	-0.007435	
2	-0.118559	0.177899	

	angletBodyGyroMeangravityMean	angletBodyGyroJerkMeangravityMean	\
0	-0.464761	-0.018446	
1	-0.732626	0.703511	
2	0.100699	0.808529	

	angleXgravityMean	angleYgravityMean	angleZgravityMean	subject	Activity	\
0	-0.841247	0.179941	-0.058627	1	5	
1	-0.844788	0.180289	-0.054317	1	5	
2	-0.848933	0.180637	-0.049118	1	5	

```

        ActivityName
0      STANDING
1      STANDING
2      STANDING

```

```
[3 rows x 564 columns]
```

```

In [4]: # get X_train and y_train from csv files
X_train = train.drop(['subject', 'Activity', 'ActivityName'], axis=1)
y_train = train.ActivityName

In [5]: # get X_test and y_test from test csv file
X_test = test.drop(['subject', 'Activity', 'ActivityName'], axis=1)
y_test = test.ActivityName

In [6]: print('X_train and y_train : ({},{})'.format(X_train.shape, y_train.shape))
        print('X_test and y_test : ({},{})'.format(X_test.shape, y_test.shape))

X_train and y_train : ((7352, 561),(7352,))
X_test and y_test : ((2947, 561),(2947,))

```

1 Let's model with our data

1.0.1 Labels that are useful in plotting confusion matrix

```
In [7]: labels=['LAYING', 'SITTING', 'STANDING', 'WALKING', 'WALKING_DOWNSTAIRS', 'WALKING_UPSTAIRS']
```

1.0.2 Function to plot the confusion matrix

```

In [8]: import itertools
import numpy as np
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix
plt.rcParams["font.family"] = 'DejaVu Sans'

def plot_confusion_matrix(cm, classes,
                          normalize=False,
                          title='Confusion matrix',
                          cmap=plt.cm.Blues):
    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]

    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation=90)

```

```

plt.yticks(tick_marks, classes)

fmt = '.2f' if normalize else 'd'
thresh = cm.max() / 2.
for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
    plt.text(j, i, format(cm[i, j], fmt),
             horizontalalignment="center",
             color="white" if cm[i, j] > thresh else "black")

plt.tight_layout()
plt.ylabel('True label')
plt.xlabel('Predicted label')

```

1.0.3 Generic function to run any model specified

```

In [9]: from datetime import datetime
def perform_model(model, X_train, y_train, X_test, y_test, class_labels, cm_normalize=
            print_cm=True, cm_cmap=plt.cm.Greens):

    # to store results at various phases
    results = dict()

    # time at which model starts training
    train_start_time = datetime.now()
    print('training the model..')
    model.fit(X_train, y_train)
    print('Done \n \n')
    train_end_time = datetime.now()
    results['training_time'] = train_end_time - train_start_time
    print('training_time(HH:MM:SS.ms) - {}\n\n'.format(results['training_time']))

    # predict test data
    print('Predicting test data')
    test_start_time = datetime.now()
    y_pred = model.predict(X_test)
    test_end_time = datetime.now()
    print('Done \n \n')
    results['testing_time'] = test_end_time - test_start_time
    print('testing time(HH:MM:SS.ms) - {}\n\n'.format(results['testing_time']))
    results['predicted'] = y_pred

    # calculate overall accuracy of the model
    accuracy = metrics.accuracy_score(y_true=y_test, y_pred=y_pred)
    # store accuracy in results
    results['accuracy'] = accuracy

```

```

print('-----')
print('|      Accuracy      |')
print('-----')
print('\n    {} \n \n'.format(accuracy))

# confusion matrix
cm = metrics.confusion_matrix(y_test, y_pred)
results['confusion_matrix'] = cm
if print_cm:
    print('-----')
    print('| Confusion Matrix |')
    print('-----')
    print('\n {} \n \n'.format(cm))

# plot confusion matrix
plt.figure(figsize=(8,8))
plt.grid(b=False)
plot_confusion_matrix(cm, classes=class_labels, normalize=True, title='Normalized Confusion Matrix')
plt.show()

# get classification report
print('-----')
print('| Classification Report |')
print('-----')
classification_report = metrics.classification_report(y_test, y_pred)
# store report in results
results['classification_report'] = classification_report
print(classification_report)

# add the trained model to the results
results['model'] = model

return results

```

1.0.4 Method to print the gridsearch Attributes

```

In [10]: def print_grid_search_attributes(model):
    # Estimator that gave highest score among all the estimators formed in GridSearch
    print('-----')
    print('|      Best Estimator      |')
    print('-----')
    print('\n\t{} \n \n'.format(model.best_estimator_))

    # parameters that gave best results while performing grid search

```

```

print('-----')
print('|      Best parameters      |')
print('-----')
print('\tParameters of best estimator : \n\n\t{}\n'.format(model.best_params_))

# number of cross validation splits
print('-----')
print('|    No of CrossValidation sets    |')
print('-----')
print('\n\tTotal numbere of cross validation sets: {}\n'.format(model.n_splits_))

# Average cross validated score of the best estimator, from the Grid Search
print('-----')
print('|      Best Score      |')
print('-----')
print('\n\tAverage Cross Validate scores of best estimator : \n\n\t{}\n'.format(model.best_score_))

```

2 1. Logistic Regression with Grid Search

```

In [11]: from sklearn import linear_model
         from sklearn import metrics

         from sklearn.model_selection import GridSearchCV

In [13]: # start Grid search
         parameters = {'C':[0.01, 0.1, 1, 10, 20, 30], 'penalty':['l2','l1']}
         log_reg = linear_model.LogisticRegression()
         log_reg_grid = GridSearchCV(log_reg, param_grid=parameters, cv=3, verbose=1, n_jobs=-1)
         log_reg_grid_results = perform_model(log_reg_grid, X_train, y_train, X_test, y_test,

```

training the model..

Fitting 3 folds for each of 12 candidates, totalling 36 fits

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.

[Parallel(n_jobs=-1)]: Done 36 out of 36 | elapsed: 1.4min finished

c:\users\dell\appdata\local\programs\python\python36\lib\site-packages\sklearn\linear_model\logistic.py:1786: FutureWarning:

c:\users\dell\appdata\local\programs\python\python36\lib\site-packages\sklearn\linear_model\logistic.py:1786: "this warning.", FutureWarning)

Done

```
training_time(HH:MM:SS.ms) - 0:01:32.121277
```

```
Predicting test data  
Done
```


```
testing_time(HH:MM:SS.ms) - 0:00:00.005913
```

```
-----  
|      Accuracy      |  
-----
```

```
0.9630132337970818
```

```
-----  
| Confusion Matrix |  
-----
```

```
[[537  0  0  0  0  0]  
[ 2 428 57  0  0  4]  
[ 0 11 520  1  0  0]  
[ 0  0  0 495  1  0]  
[ 0  0  0  3 409  8]  
[ 0  0  0 22  0 449]]
```




Classification Report

	precision	recall	f1-score	support
LAYING	1.00	1.00	1.00	537
SITTING	0.97	0.87	0.92	491
STANDING	0.90	0.98	0.94	532
WALKING	0.95	1.00	0.97	496
WALKING_DOWNSTAIRS	1.00	0.97	0.99	420
WALKING_UPSTAIRS	0.97	0.95	0.96	471
micro avg	0.96	0.96	0.96	2947

macro avg	0.97	0.96	0.96	2947
weighted avg	0.96	0.96	0.96	2947

```
In [14]: plt.figure(figsize=(8,8))
plt.grid(b=False)
plot_confusion_matrix(log_reg_grid_results['confusion_matrix'], classes=labels, cmap=)
plt.show()
```



```
In [15]: # observe the attributes of the model
print_grid_search_attributes(log_reg_grid_results['model'])
```

```
-----
| Best Estimator |
```

```
LogisticRegression(C=30, class_weight=None, dual=False, fit_intercept=True,
    intercept_scaling=1, max_iter=100, multi_class='warn',
    n_jobs=None, penalty='l2', random_state=None, solver='warn',
    tol=0.0001, verbose=0, warm_start=False)
```

| Best parameters |

Parameters of best estimator :

```
{'C': 30, 'penalty': 'l2'}
```

| No of CrossValidation sets |

Total nombre of cross validation sets: 3

| Best Score |

Average Cross Validate scores of best estimator :

0.9458650707290533

3 2. Linear SVC with GridSearch

```
In [16]: from sklearn.svm import LinearSVC
```

```
In [17]: parameters = {'C':[0.125, 0.5, 1, 2, 8, 16]}
         lr_svc = LinearSVC(tol=0.00005)
         lr_svc_grid = GridSearchCV(lr_svc, param_grid=parameters, n_jobs=-1, verbose=1)
         lr_svc_grid_results = perform_model(lr_svc_grid, X_train, y_train, X_test, y_test, cl
```

training the model..

Fitting 3 folds for each of 6 candidates, totalling 18 fits

```
c:\users\dell\appdata\local\programs\python\python36\lib\site-packages\sklearn\model_selection\
warnings.warn(CV_WARNING, FutureWarning)
```

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
```

```
[Parallel(n_jobs=-1)]: Done 18 out of 18 | elapsed: 26.6s finished
```

```
c:\users\dell\appdata\local\programs\python\python36\lib\site-packages\sklearn\svm\base.py:922
```

```
"the number of iterations.", ConvergenceWarning)
```

```
Done
```

```
training_time(HH:MM:SS.ms) - 0:00:30.736681
```

```
Predicting test data
```

```
Done
```


```
testing time(HH:MM:SS.ms) - 0:00:00.006907
```

```
-----  
|      Accuracy      |  
-----
```

```
0.9657278588394977
```

```
-----  
| Confusion Matrix |  
-----
```

```
[[537  0  0  0  0  0]  
[ 1 423 64  0  0  3]  
[ 0  9 523  0  0  0]  
[ 0  0  0 496  0  0]  
[ 0  0  0  2 413  5]  
[ 0  0  0 16  1 454]]
```



Classification Report

	precision	recall	f1-score	support
LAYING	1.00	1.00	1.00	537
SITTING	0.98	0.86	0.92	491
STANDING	0.89	0.98	0.93	532
WALKING	0.96	1.00	0.98	496
WALKING_DOWNSTAIRS	1.00	0.98	0.99	420
WALKING_UPSTAIRS	0.98	0.96	0.97	471
micro avg	0.97	0.97	0.97	2947

macro avg	0.97	0.97	0.97	2947
weighted avg	0.97	0.97	0.97	2947

```
In [18]: print_grid_search_attributes(lr_svc_grid_results['model'])
```

```
-----
|      Best Estimator      |
|-----|
```

```
LinearSVC(C=2, class_weight=None, dual=True, fit_intercept=True,
intercept_scaling=1, loss='squared_hinge', max_iter=1000,
multi_class='ovr', penalty='l2', random_state=None, tol=5e-05,
verbose=0)
```

```
-----
|    Best parameters      |
|-----|
```

```
Parameters of best estimator :
```

```
{'C': 2}
```

```
-----
| No of CrossValidation sets |
|-----|
```

```
Total nombre of cross validation sets: 3
```

```
-----
|      Best Score      |
|-----|
```

```
Average Cross Validate scores of best estimator :
```

```
0.9469532100108814
```

4 3. Kernel SVM with GridSearch

```
In [19]: from sklearn.svm import SVC
parameters = {'C':[2,8,16],\
              'gamma': [ 0.0078125, 0.125, 2]}
rbf_svm = SVC(kernel='rbf')
rbf_svm_grid = GridSearchCV(rbf_svm,param_grid=parameters, n_jobs=-1)
rbf_svm_grid_results = perform_model(rbf_svm_grid, X_train, y_train, X_test, y_test, c
```

training the model..

```
c:\users\dell\appdata\local\programs\python\python36\lib\site-packages\sklearn\model_selection\
warnings.warn(CV_WARNING, FutureWarning)
```

Done

training_time(HH:MM:SS.ms) - 0:05:48.772455

Predicting test data

Done


testing time(HH:MM:SS.ms) - 0:00:02.371631

```
-----
|      Accuracy      |
-----
```

0.9626739056667798

```
-----
| Confusion Matrix |
-----
```

```
[[537  0  0  0  0  0]
 [ 0 441 48  0  0  2]
 [ 0 12 520  0  0  0]
 [ 0  0  0 489  2  5]
 [ 0  0  0  4 397 19]
 [ 0  0  0 17  1 453]]
```



Classification Report

	precision	recall	f1-score	support
LAYING	1.00	1.00	1.00	537
SITTING	0.97	0.90	0.93	491
STANDING	0.92	0.98	0.95	532
WALKING	0.96	0.99	0.97	496
WALKING_DOWNSTAIRS	0.99	0.95	0.97	420
WALKING_UPSTAIRS	0.95	0.96	0.95	471
micro avg	0.96	0.96	0.96	2947

macro avg	0.96	0.96	0.96	2947
weighted avg	0.96	0.96	0.96	2947

```
In [20]: print_grid_search_attributes(rbf_svm_grid_results['model'])
```

```
-----
|           Best Estimator           |
|-----|
```

```
SVC(C=16, cache_size=200, class_weight=None, coef0=0.0,
decision_function_shape='ovr', degree=3, gamma=0.0078125, kernel='rbf',
max_iter=-1, probability=False, random_state=None, shrinking=True,
tol=0.001, verbose=False)
```

```
-----
|           Best parameters           |
|-----|
```

```
Parameters of best estimator :
```

```
{'C': 16, 'gamma': 0.0078125}
```

```
-----
| No of CrossValidation sets         |
|-----|
```

```
Total nombre of cross validation sets: 3
```

```
-----
|           Best Score               |
|-----|
```

```
Average Cross Validate scores of best estimator :
```

```
0.9440968443960827
```

5 4. Decision Trees with GridSearchCV

```
In [21]: from sklearn.tree import DecisionTreeClassifier
parameters = {'max_depth':np.arange(3,10,2)}
dt = DecisionTreeClassifier()
dt_grid = GridSearchCV(dt,param_grid=parameters, n_jobs=-1)
dt_grid_results = perform_model(dt_grid, X_train, y_train, X_test, y_test, class_label)
print_grid_search_attributes(dt_grid_results['model'])
```

training the model..

```
c:\users\dell\appdata\local\programs\python\python36\lib\site-packages\sklearn\model_selection\
warnings.warn(CV_WARNING, FutureWarning)
```

Done

training_time(HH:MM:SS.ms) - 0:00:08.387619

Predicting test data

Done


testing time(HH:MM:SS.ms) - 0:00:00.005944

```
-----
|      Accuracy      |
-----
```

0.8635900916185952

```
-----
| Confusion Matrix |
-----
```

```
[[537  0  0  0  0  0]
 [ 0 386 105  0  0  0]
 [ 0  93 439  0  0  0]
 [ 0  0  0 472 16  8]
 [ 0  0  0 17 342 61]
 [ 0  0  0 78 24 369]]
```

Classification Report

	precision	recall	f1-score	support
LAYING	1.00	1.00	1.00	537
SITTING	0.81	0.79	0.80	491
STANDING	0.81	0.83	0.82	532
WALKING	0.83	0.95	0.89	496
WALKING_DOWNSTAIRS	0.90	0.81	0.85	420
WALKING_UPSTAIRS	0.84	0.78	0.81	471
micro avg	0.86	0.86	0.86	2947

macro avg	0.86	0.86	0.86	2947
weighted avg	0.86	0.86	0.86	2947

```
-----
|      Best Estimator      |
|-----|
```

```
DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=7,
    max_features=None, max_leaf_nodes=None,
    min_impurity_decrease=0.0, min_impurity_split=None,
    min_samples_leaf=1, min_samples_split=2,
    min_weight_fraction_leaf=0.0, presort=False, random_state=None,
    splitter='best')
```

```
-----
|    Best parameters      |
|-----|
```

Parameters of best estimator :

```
{'max_depth': 7}
```

```
-----
|  No of CrossValidation sets  |
|-----|
```

Total nombre of cross validation sets: 3

```
-----
|      Best Score        |
|-----|
```

Average Cross Validate scores of best estimator :

```
0.8392274211099021
```

6 5. Random Forest Classifier with GridSearch

```
In [22]: from sklearn.ensemble import RandomForestClassifier
        params = {'n_estimators': np.arange(10,201,20), 'max_depth':np.arange(3,15,2)}
        rfc = RandomForestClassifier()
        rfc_grid = GridSearchCV(rfc, param_grid=params, n_jobs=-1)
        rfc_grid_results = perform_model(rfc_grid, X_train, y_train, X_test, y_test, class_label)
        print_grid_search_attributes(rfc_grid_results['model'])
```

training the model..

```
c:\users\dell\appdata\local\programs\python\python36\lib\site-packages\sklearn\model_selection\
warnings.warn(CV_WARNING, FutureWarning)
```

Done

```
training_time(HH:MM:SS.ms) - 0:03:15.851561
```

Predicting test data

Done


```
testing time(HH:MM:SS.ms) - 0:00:00.046874
```

```
-----
|      Accuracy      |
|-----|
```

0.9158466236851035

```
-----
| Confusion Matrix |
|-----|
```

```
[[537  0  0  0  0  0]
 [ 0 421 70  0  0  0]
 [ 0 38 494  0  0  0]
 [ 0  0  0 485  9  2]
 [ 0  0  0 35 337 48]
 [ 0  0  0 40  6 425]]
```



| Classification Report |

	precision	recall	f1-score	support
LAYING	1.00	1.00	1.00	537
SITTING	0.92	0.86	0.89	491
STANDING	0.88	0.93	0.90	532
WALKING	0.87	0.98	0.92	496
WALKING_DOWNSTAIRS	0.96	0.80	0.87	420
WALKING_UPSTAIRS	0.89	0.90	0.90	471
micro avg	0.92	0.92	0.92	2947

macro avg	0.92	0.91	0.91	2947
weighted avg	0.92	0.92	0.92	2947

```
-----
|      Best Estimator      |
|-----|
```

```
RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
    max_depth=7, max_features='auto', max_leaf_nodes=None,
    min_impurity_decrease=0.0, min_impurity_split=None,
    min_samples_leaf=1, min_samples_split=2,
    min_weight_fraction_leaf=0.0, n_estimators=130, n_jobs=None,
    oob_score=False, random_state=None, verbose=0,
    warm_start=False)
```

```
-----
|      Best parameters      |
|-----|
```

Parameters of best estimator :

```
{'max_depth': 7, 'n_estimators': 130}
```

```
-----
|  No of CrossValidation sets  |
|-----|
```

Total nombre of cross validation sets: 3

```
-----
|      Best Score      |
|-----|
```

Average Cross Validate scores of best estimator :

```
0.9151251360174102
```

7 6. Gradient Boosted Decision Trees With GridSearch

```
In [23]: from sklearn.ensemble import GradientBoostingClassifier
        param_grid = {'max_depth': np.arange(5,8,1), \
            'n_estimators':np.arange(130,170,10)}
        gbdt = GradientBoostingClassifier()
        gbdt_grid = GridSearchCV(gbdt, param_grid=param_grid, n_jobs=-1)
        gbdt_grid_results = perform_model(gbdt_grid, X_train, y_train, X_test, y_test, class_)
        print_grid_search_attributes(gbdt_grid_results['model'])
```

training the model..

```
c:\users\dell\appdata\local\programs\python\python36\lib\site-packages\sklearn\model_selection\
warnings.warn(CV_WARNING, FutureWarning)
```

Done

training_time(HH:MM:SS.ms) - 0:28:30.458286

Predicting test data

Done


testing time(HH:MM:SS.ms) - 0:00:00.077851

```
-----
|      Accuracy      |
-----
```

0.9229725144214456

```
-----
| Confusion Matrix |
-----
```

```
[[537  0  0  0  0  0]
 [ 0 398 91  0  0  2]
 [ 0 37 495  0  0  0]
 [ 0  0  0 483  7  6]
 [ 0  0  0 10 374 36]
 [ 0  1  0 31  6 433]]
```



Classification Report

	precision	recall	f1-score	support
LAYING	1.00	1.00	1.00	537
SITTING	0.91	0.81	0.86	491
STANDING	0.84	0.93	0.89	532
WALKING	0.92	0.97	0.95	496
WALKING_DOWNSTAIRS	0.97	0.89	0.93	420
WALKING_UPSTAIRS	0.91	0.92	0.91	471
micro avg	0.92	0.92	0.92	2947

macro avg	0.93	0.92	0.92	2947
weighted avg	0.92	0.92	0.92	2947

```
-----
|      Best Estimator      |
|-----|
```

```
GradientBoostingClassifier(criterion='friedman_mse', init=None,
    learning_rate=0.1, loss='deviance', max_depth=5,
    max_features=None, max_leaf_nodes=None,
    min_impurity_decrease=0.0, min_impurity_split=None,
    min_samples_leaf=1, min_samples_split=2,
    min_weight_fraction_leaf=0.0, n_estimators=160,
    n_iter_no_change=None, presort='auto', random_state=None,
    subsample=1.0, tol=0.0001, validation_fraction=0.1,
    verbose=0, warm_start=False)
```

```
-----
|      Best parameters      |
|-----|
```

Parameters of best estimator :

```
{'max_depth': 5, 'n_estimators': 160}
```

```
-----
|  No of CrossValidation sets  |
|-----|
```

Total numbere of cross validation sets: 3

```
-----
|      Best Score      |
|-----|
```

Average Cross Validate scores of best estimator :

```
0.904923830250272
```

8 7. Comparing all models

```
In [24]: print('\n
          print('
          print('Logistic Regression : {:.04}%
                                     Accuracy      Error')
                                     -----      -----')
                                     {:.04}%      {:.04}%'.format(log_reg_grid_results['accuracy'],
                                     100-(log_reg_grid_results['accuracy']
```



```

print('Linear SVC          : {:.04}%      {:.04}% '.format(lr_svc_grid_results['accuracy'],
                                                             100-(lr_svc_grid_results['accuracy'])))

print('rbf SVM classifier  : {:.04}%      {:.04}% '.format(rbf_svm_grid_results['accuracy'],
                                                             100-(rbf_svm_grid_results['accuracy'])))

print('DecisionTree       : {:.04}%      {:.04}% '.format(dt_grid_results['accuracy'],
                                                             100-(dt_grid_results['accuracy'])))

print('Random Forest      : {:.04}%      {:.04}% '.format(rfc_grid_results['accuracy'],
                                                             100-(rfc_grid_results['accuracy'])))

print('GradientBoosting DT : {:.04}%      {:.04}% '.format(rfc_grid_results['accuracy'],
                                                             100-(rfc_grid_results['accuracy'])))

```

	Accuracy -----	Error -----
Logistic Regression :	96.3%	3.699%
Linear SVC :	96.57%	3.427%
rbf SVM classifier :	96.27%	3.733%
DecisionTree :	86.36%	13.64%
Random Forest :	91.58%	8.415%
GradientBoosting DT :	91.58%	8.415%

We can choose *Logistic regression* or *Linear SVC* or *rbf SVM*.

9 Conclusion :

In the real world, domain-knowledge, EDA and feature-engineering matter most.